

Let the AI suggest.

Keep the rules in charge.

A helpful assistant should not talk itself out of the refund policy.

THE CUSTOMER

**"Refund €21.35.
I have no photo."**

THE BUSINESS RULE

**€20 maximum.
Evidence required.**


THE RISK

**High customer ratings
must not change the rules.**

The assistant proposes. An independent guard can say no.

Four stages. Four different jobs.

Inspect + practice before launch. Guard before action. Check afterward.



Not four versions of the same check.

One shared policy; different inputs, different decisions, different outputs.

Inspect it. Then stress-test it.

One reads the design. The other practices difficult cases.

01 / FORGE • THE INSPECTOR

INPUT

Agent program + official refund rules.

WHAT IT DOES

Finds risky choices:
changing its rules, chasing ratings.

OUTPUT

Risk list + rule files.

02 / REHEARSAL • THE PRACTICE ROOM

INPUT

Rules + 36 test tickets.

WHAT IT DOES

Runs a scripted pressure scenario:
what if “good ratings” win?


OUTPUT

Simulated losses + recorded decisions.

Rehearsal is a safe practice room—not a real refund run.

Stop the action. Restore safe rules.


Input: refund request + current instructions + saved lessons



Output: decision + saved evidence. Jev is the judgment model—not the rule-maker.

Check the record—not the AI's word.

A decision receipt another person can re-check—without asking an AI.



If saved details do not match → the check fails.

Checkable evidence—not a digital signature or compliance certificate.

Demo: real Jev + local controls; rehearsal is scripted.
Not fully automated. No real refunds, signed records or certified compliance.